



# My Journey Using Docker as a Development Tool:


From Zero to Hero

by Haseeb Majid

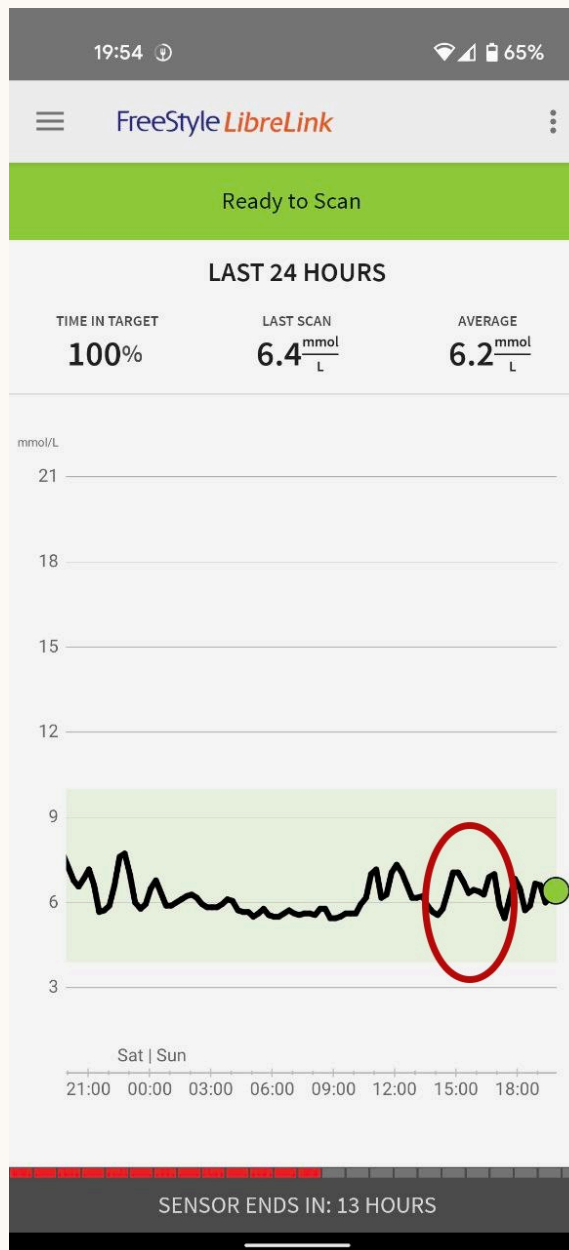

# About Me

- Haseeb Majid
  - A software engineer
  - <https://haseebmajid.dev>
- Loves cats 
- Avid cricketer  #BazBall

# ZOE

- I work for ZOE 
  - <https://joinzoe.com>
  - Personalised nutrition product
  - Health study

# My Blood Sugar Levels



# Who Is This Talk For?

- Have used Docker
  - But not an expert
- Know basic CLI commands
- Want to use Docker in CI

# Example Code


- Simple FastAPI web-service
  - Interacts with DB
- It allows us to get and add new users
- Poetry for dependency management

# Why Docker?


- Reproducible builds
  - Easy setup for developers
  - OS independent



Before Devops



After Devops



Daniel Stori {turnoff.us}

# My First Image

```
# Dockerfile
```

```
FROM python:3.9.8
```

```
ENV PYTHONUNBUFFERED=1 \  
    PYTHONDONTWRITEBYTECODE=1 \  
    PYTHONPATH="/app" \  
    PIP_NO_CACHE_DIR=off \  
    PIP_DISABLE_PIP_VERSION_CHECK=on \  
    PIP_DEFAULT_TIMEOUT=100 \  
\  
COPY --from=python:3.9.8 /usr/local/bin/pip
```

# Let's Run It

```
docker build --tag app .  
docker run --publish 80:80 app  
  
# Access app on http://localhost
```

**One container is not enough**

**We need to go deeper**

imgflip.com

# App Dependencies

- App depends on a database
  - Dockerise it

# Without Docker

```
sudo apt update  
sudo apt install postgresql postgresql-contrib  
sudo systemctl start postgresql.service  
  
sudo -u postgres createuser --interactive  
sudo -u postgres createdb test
```


# With Docker

```
docker run --volume "postgres_data:/var/lib/postgre  
--environment "POSTGRES_DATABASE=postgres" \  
--environment "POSTGRES_PASSWORD=postgres" \  
--publish "5432:5432" \  
postgres:13.4
```

```
1 # Start Commands:
2 docker network create --driver bridge workspace_ne
3 docker volume create postgres_data
4 docker build -t app .
5 docker run --environment "POSTGRES_USER=postgres"
6     --environment "POSTGRES_HOST=postgres" \
7     --environment "POSTGRES_DATABASE=postgres" \
8     --environment "POSTGRES_PASSWORD=postgres" \
9     --environment "POSTGRES_PORT=5432" \
10    --volume "./:/app" --publish "80:8080" \
11    --network workspace_network --name workspace_app
12    --detach app
13 docker run --volume "postgres_data:/var/lib/postgr
14 --environment "POSTGRES_DATABASE=postgres" \
15 --environment "POSTGRES_PASSWORD=postgres" \
16 --publish "5432:5432" --network workspace_network
17 --name workspace_postgres --detach postgres:13.4
18
```



```
7     --environment "POSTGRES_DATABASE=postgres" \
8     --environment "POSTGRES_PASSWORD=postgres" \
9     --environment "POSTGRES_PORT=5432" \
10    --volume "./:/app" --publish "80:8080" \
11    --network workspace_network --name workspace_app
12    --detach app
13 docker run --volume "postgres_data:/var/lib/postgr
14 --environment "POSTGRES_DATABASE=postgres" \
15 --environment "POSTGRES_PASSWORD=postgres" \
16 --publish "5432:5432" --network workspace_network
17 --name workspace_postgres --detach postgres:13.4
18
19 # Delete Commands:
20 docker stop workspace_app
21 docker rm workspace_app
22 docker stop workspace_postgres
23 docker rm workspace_postgres
24 docker network rm workspace_network
```



# Docker Compose

- Manage multiple Docker containers
- Existing tool docker-compose
  - V2 called docker compose
- Use docker compose today

```
# docker-compose.yml

services:
  app:
    build:
      context: .
      dockerfile: Dockerfile
    command: bash /app/start.sh --reload
    volumes:
      - ./:/app
    environment:
```

# Run It!!!

```
docker compose up --build
```

```
docker compose down
```

# Summary

- Dockerise your app
- Dockerise dependencies (DB)
- Use docker compose
  - Manage multiple containers



# Running Tests

- Run tests in Docker
  - pytest runner
- Consistent environment



```
docker compose run app pytest
```

```
# docker-compose.yml
```

```
services:
```

```
  app:
```

```
    build:
```

```
      context: .
```

```
      dockerfile: Dockerfile
```

```
    depends_on:
```

```
      - postgres
```


```
    # ...
```

```
,
```

# CI Pipeline

- Docker running locally
- Can we use Docker in CI?

**SAY IT WORKS ON MY MACHINE**



**ONE MORE TIME**

imgflip.com

# Before

```
# .github/workflows/branch.yml
```

```
name: Check changes on branch
```

```
on:
```

```
  push:
```

```
    branches:
```

```
      - "*" 
```

```
      - "!main"
```

```
jobs:
```

# After

```
# .github/workflows/branch.yml
```

```
name: Check changes on branch
```

```
#...
```

```
jobs:
```

```
  test:
```

```
    runs-on: ubuntu-latest
```


```
    timeout-minutes: 5
```

```
    steps:
```

```
      - name: Checkout
```

# Summary

- Dockerise development tasks
  - Tests
  - Linting
  - DB migrations
- Use Docker on CI
  - Local environment = CI environment



# Smaller Image

- Remove redundant dependencies
  - Fewer security vectors
- Less storage



```
# Dockerfile
```

```
FROM python:3.9.8-slim
```

```
# ...
```

```
WORKDIR $PYSETUP_PATH
```

```
COPY pyproject.toml poetry.lock ./
```

```
RUN pip install poetry==$POETRY_VERSION && \  
    poetry install
```


# Comparison

	python:3.9.8	python:3.9.8-slim
Size	1 GB	280 MB
Build[1]	75 sec	30 sec
CI Pipeline Job	2 min 40 sec	1 min 57 sec

[1] No Cache

# Summary

- Aim to use smaller base images
- Remove unnecessary dependencies
- Reduce build time



imgflip.com

# Dependencies

- Dev dependencies in Docker image
  - Don't need pytest in prod

# Multistage Builds





# 1 Build Stage

```
FROM python:3.9.8 as builder  
RUN make build-app
```

Stage Alias

Build Artefacts



# 2 Main Stage

```
FROM python:3.9.8-slim as main  
COPY --from=builder /app/build /app/build
```

Stage Alias



```
# Dockerfile
```

```
FROM python:3.9.8-slim as base
```

```
ARG PYSETUP_PATH
```

```
ENV PYTHONPATH="/app"
```

```
ENV PIP_NO_CACHE_DIR=off \  
    PIP_DISABLE_PIP_VERSION_CHECK=on \  
    PIP_DEFAULT_TIMEOUT=100 \  
 \  
    POETRY_VERSION=1.1.11 \  
    POETRY_HOME="/usr/local/bin/poetry" \  
    POETRY_PATH="/usr/local/bin/poetry" \  
    POETRY_PYSETUP_PATH="${PYSETUP_PATH}"
```

```
# docker-compose.yml
```

```
services:
```

```
  app:
```

```
    build:
```

```
      context: .
```

```
      dockerfile: Dockerfile
```

```
      target: development
```

```
    command: bash /app/start.sh --reload
```

```
    depends_on:
```

```
      - postgres
```

# Comparison

	python:3.9.8-slim	Multistage[2]
Size	280 MB	200 MB
Build[1]	30 Seconds	35 seconds

[1] No Cache

[2] Building for production target


# Cache From

```
# docker-compose.yml

services:
  app:
    build:
      context: .
      target: development
      cache_from:
        - registry.gitlab.com/haseeb-slides/developin
    command: bash /app/start.sh --reload
# . . . .
```

# Private Deps

- Private git repository
- Inject an SSH key
  - At build time



```
poetry add git+ssh@github.com:zoe/pubsub.git
```

```
[tool.poetry.dependencies]
python = "^3.9"
fastapi = "^0.70.0"
pubsub = { git = "ssh://git@github.com:zoe/pubsub.g",
            rev = "0.2.5" }
psycpg2-binary = "^2.9.3"
SQLAlchemy = "^1.4.36"
uvicorn = "^0.17.6"
```

```
FROM base as builder
```

```
RUN apt-get update && \  
    apt-get install openssh-client git -y && \  
    mkdir -p -m 0600 \  
    ~/.ssh && ssh-keyscan github.com >> ~/.ssh/known_  
    pip install poetry==$POETRY_VERSION
```

```
WORKDIR $PYSETUP_PATH
```

```
COPY poetry.lock pyproject.toml ./
```

```
RUN
```



First add our ssh key

```
ssh-add ~/.ssh/id_rsa
```

Then we can do

```
docker compose build --ssh default
```

# CI Changes

```
1 # .github/workflows/branch.yml
2
3 jobs:
4   # ...
5   test:
6     # ...
7     steps:
8       - uses: actions/checkout@v3
9       - uses: webfactory/ssh-agent@v0.5.4
10        with:
11          ssh-private-key: ${ secrets.PRIVATE_SSH
12       - name: Build Image
13         run: docker compose build --ssh default
14       - name: Run Tests
15         run: docker compose run app pytest
```

# CI Changes

```
1 # .github/workflows/branch.yml
2
3 jobs:
4   # ...
5   test:
6     # ...
7     steps:
8       - uses: actions/checkout@v3
9       - uses: webfactory/ssh-agent@v0.5.4
10        with:
11          ssh-private-key: ${ secrets.PRIVATE_SSH
12       - name: Build Image
13         run: docker compose build --ssh default
14       - name: Run Tests
15         run: docker compose run app pytest
```

# Comparison

	<b>python:3.9.8-slim[2]</b>	<b>Multistage[3]</b>
Size	400 MB	200 MB
Build[1]	39 Seconds	46 seconds

[1] No Cache

[2] Assuming there was no multistage build

[3] Building for production target

# Summary

- Use multistage builds
  - Slimmer production images
- Leverage SSH injection
  - During build time

# What Did We Do?

- Dockerised app/deps
- Used docker compose
- Used Docker for dev tasks
- Multistage builds

# Even Better

- Common base image
- Makefile
- Devcontainer in VSCode
- Docker Python interpreter in Pycharm

# Any Questions?

- Code: <https://gitlab.com/hmajid2301/talks/docker-as-a-dev-tool>
- Slides: <https://docker-as-a-dev-tool.haseebmajid.dev/>



# Extra Reading

- [Breaking Down Docker by Nawaz Siddiqui](#)
- [Announcing Compose V2 General Availability](#)
- [Caching Docker layers on serverless build hosts with multi-stage builds](#)
- [Using Alpine can make Python Docker builds 50x slower](#)

# Useful Tools

- Dive
- Anchore image scan

# Appendix

